

Team members

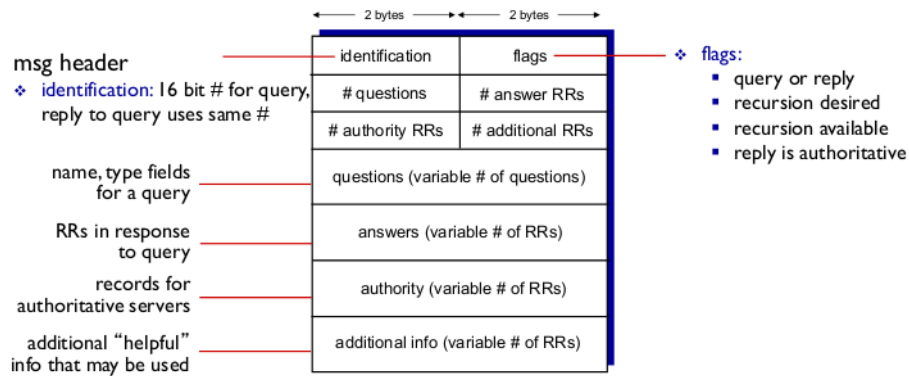
Name	Roll no.
Pranav Joshi	22110197
Nishit	22110172

Task 1: Custom DNS Resolver

Theory

DNS

The Domain Name System is a distributed database service that maps hostnames (www.google.com for example) to IP addresses (142.251.220.35 for example). The message format for DNS is as follows :



The DNS message payload can be extracted using **pyshark** which is a Python wrapper around TShark, the CLI alternative to the WireShark software. Since for this task, we only need to extract the queries, we need to look at bytes 5 and 6 (with 1 based indexing) to get the number of queries/questions, and then extract the questions, starting from byte 13.

For example, this is an DNS message (in hexadecimal, and bytes separated by ':') :

00:00:01:00:00:01:00:00:00:00:00:00:07:74:77:69:74:74:65:72:03:63:6f:6d:00:00:01:00:01

Here, we have :

Field	Value (hexadecimal)
Identification	00 00
Flags	01 00
no. of questions	00 01

Field	Value (hexadecimal)
no. of answers	00 00
no. of authority RRs	00 00
no. of additional RRs	00 00
questions	07 74 77 69 74 74 65 72 03 63 6f 6d 00 00 01 00 01

The number of question is thus 1. In the “questions” section, each question is a (name,type,class) tuple. The name part is encoded by breaking the Fully Qualified Domain Name (FQDN) into its constituent parts (for example, “twitter.com” to [“twitter”,“com”]) and then encoding each character in ASCII format (one byte for each character) for each part. For each part, we also add prepend a byte that tells the length. For example, for “twitter” the length is 7 and so we prepend the byte 0x07 before the encoded name, namely 74 77 69 74 74 65 72 (hexadecimal). Similarly, for “com” , we get length of 0x03 and encoding of 63 6f 6d . Finally, we put a null character 0x00 signifying the end of the name. After that, the next 2 bytes are the type. In our example, it is 0x0001 (1 in decimal), which is the code for A type queries. And then, the last 2 bytes are the class . In our example, it is 0x0001, which is the code for IN (means “internet”).

Since we are asked to only report the custom header, FQDN, and the response from the server, other information, such as type and class are not needed to be extracted.

Rather than parsing the bytes myself, I will rely on **pyshark** to do this by default.

But just for the sake of completeness, I’ll provide an option (**--parse**) that will parse the DNS message manually and extract the queries.

Custom header

According to the assignment PDF, the custom header is

A value of the current date-time and id, use timestamp in format “HHMMSSID”

This can be interpreted as either the current system time on the client, or the time when the DNS message was originally sent. The time when DNS message is sent is recorded by WireShark while capturing the packets and can be easily extracted by **pyshark**. Since I am unsure what to implement, I will use the current system time on my laptop by default and provide the **--st** option to use the packet sniff time.

File Structure

- **client.py**: Client-side script for packet parsing and message sending.
- **server.py**: Server-side script for DNS resolution logic.

- **DNSmessages.csv**: Stores extracted DNS queries. This is optional and can be avoided using `--dwem` option for `client.py`.
- **Report.csv**: Records query headers, domains, and resolved IPs.
- **9.pcap**: The PCAP file number should be $(197 + 172) \bmod 10 = 9$. Thus, I should use **9.pcap** for the task.

Usage

Prerequisites

- Python 3.x
- pyshark library (`pip install pyshark`)
- PCAP file for DNS traffic (see assignment instructions for correct selection)

Running the Server

```
python3 server.py
```

This will start the server on `127.0.0.1:53535`, giving this output:

```
Server listening on 127.0.0.1:53535
```

Ensure the server is running before starting the client.

Running the Client

```
python3 client.py [OPTIONS] <pcapfiles>
```

You can pass in multiple (space separated) PCAP filenames in place of `<pcapfiles>`.

Options:

- `--help`: Show help message and usage instructions.
- `--de`: Skip extraction from PCAP and use previously extracted messages (in **DNSmessages.csv**).
- `--dwem`: Skip writing extracted messages in **DNSmessages.csv**.
- `--tcp`: Uses TCP rather than UDP for transport layer.
- `--parse`: Parses the DNS message without relying on PyShark
- `--st`: Uses packet sniff time rather than the current system time.
- `--batch`: Batches the extracted DNS messages to resolve together, rather than immediately after extraction (default)

The client parses DNS queries from the PCAP file, adds a header, sends them to the server, and writes the results to **Report.csv**.

Running the server

```
python3 server.py [--tcp]
```

If the `--tcp` option is set, the server uses TCP to communicate with the client. When using UDP, there is a thread that waits for the `exit` command on the terminal to gracefully close the server. Pressing `Ctrl+C` also exits (for both UDP and TCP modes).

Example

On client side we get this:

```
$ python3 client.py 9.pcap
Packet : 2025-09-11 13:50:22.885383 : UDP 10.240.26.55 → 8.8.8.8 : 00:00:01:0...
Sent 37 bytes to 127.0.0.1:53535
Response: 192.168.1.6
Packet : 2025-09-11 13:50:33.896118 : UDP 10.240.26.55 → 8.8.8.8 : 00:00:01:0...
Sent 37 bytes to 127.0.0.1:53535
Response: 192.168.1.7
Packet : 2025-09-11 13:50:37.059897 : UDP 10.240.26.55 → 8.8.8.8 : 00:00:01:0...
Sent 37 bytes to 127.0.0.1:53535
Response: 192.168.1.8
Packet : 2025-09-11 13:50:45.720688 : UDP 10.240.26.55 → 8.8.8.8 : 00:00:01:0...
Sent 38 bytes to 127.0.0.1:53535
Response: 192.168.1.9
Packet : 2025-09-11 13:50:53.805059 : UDP 10.240.26.55 → 8.8.8.8 : 00:00:01:0...
Sent 36 bytes to 127.0.0.1:53535
Response: 192.168.1.10
Packet : 2025-09-11 13:51:03.031684 : UDP 10.240.26.55 → 8.8.8.8 : 00:00:01:0...
Sent 36 bytes to 127.0.0.1:53535
Response: 192.168.1.6
Extracted DNS queries from PCAP and resolved
```

and on server side :

```
$ python3 server.py
Server listening on 127.0.0.1:53535
Received DNS query from 127.0.0.1:43863
Header extracted :13:50:22:0
IP to return : 192.168.1.6
response sent
Received DNS query from 127.0.0.1:59360
Header extracted :13:50:33:1
IP to return : 192.168.1.7
response sent
Received DNS query from 127.0.0.1:47293
Header extracted :13:50:37:2
IP to return : 192.168.1.8
response sent
Received DNS query from 127.0.0.1:36421
```

```

Header extracted :13:50:45:3
IP to return : 192.168.1.9
response sent
Received DNS query from 127.0.0.1:39677
Header extracted :13:50:53:4
IP to return : 192.168.1.10
response sent
Received DNS query from 127.0.0.1:43491
Header extracted :13:51:03:5
IP to return : 192.168.1.6
response sent
exit
Shutting down server...

```

The output (in `Report.csv`) will be:

Custom header value (HHMMSSID)	Domain name	Resolved IP address
13502200	twitter.com	192.168.1.6
13503301	example.com	192.168.1.7
13503702	netflix.com	192.168.1.8
13504503	linkedin.com	192.168.1.9
13505304	reddit.com	192.168.1.10
13510305	openai.com	192.168.1.6

Output Files

- `DNSmessages.csv`: Extracted DNS queries from the PCAP (for inspection).
- `Report.csv`: Table of each query, header, and resolved IP (for submission).

Task 2: Traceroute on different OS

Theory

traceroute

According to the manual page for **traceroute** 2.1.0, the way this utility measures the Round-Trip-Time (RTT) for each node in the path to a particular destination host is by sending probing packets (with various transport and network layer protocols) with small time-to-live (TTL) values. It starts with a TTL value of 1 and increases upto `max_ttl` which is set to 30 by default. For every TTL value, it sends 3 probe packets. For every packet that exceeds the number of hops (the TTL value) while moving towards the destination, the router which last decremented its TTL value (before it reaching 0) sends back an ICMP packet that reports this event. The time taken for this packet to reach, starting from the time that the probe packet was sent is the RTT for that node, whose IP address is there in said ICMP packet.

tcpdump

Now, **tcpdump** is an utility that allows capturing the packets sent to your device. Since I am connected to the internet through an ethernet connection with my smartphone (using USB), which is further connected to the internet using cellular data, my laptop won't be getting any other traffic. Now, I will run **tcpdump** in th background as `sudo tcpdump -i enx0207646b3637 -w traceroutecap.pcap & .` Here, `enx0207646b3637` is the USB ethernet interface and `traceroutecap.pcap` is the name of the output file where the capture/dump will be written. To kill this **tcpdump** process later, I will store the process ID as `TCPDUMP_PID=$!`. Then, I will run **traceroute** and later kill the **tcpdump** process using `sudo kill $TCPDUMP_PID .` Note that my smartphone's IPv4 address is 49.34.205.233 and my laptop's IPv4 address is 192.168.42.179 . This will come in handy later. This modification is harmless to our analysis as I'll show later (after answering the questions).

Execution

What I really want to run is **traceroute google.com**, but I don't want to deal with the packets involved in the DNS lookup for **google.com**. So, I will instead, use the IPv4 address 142.250.183.78. This is one of Google's servers.

I got the IPv4 address using **nslookup**.

```
~$ nslookup google.com
Server:      127.0.0.53
Address:     127.0.0.53#53
```

```
Non-authoritative answer:
Name:   google.com
```

Address: 142.250.183.78
Name: google.com
Address: 2404:6800:4009:821::200e

Note that Google uses load balancing and the output for `nslookup` is not constant.
The final script for the capture is :

```
sudo tcpdump -i enx0207646b3637 -w traceroutecap.pcap &  
TCPDUMP_PID=$!  
traceroute 142.250.183.78  
sudo kill $TCPDUMP_PID
```

The output of `traceroute` is :

```
traceroute to 142.250.183.78 (142.250.183.78), 30 hops max, 60 byte packets  
 1  _gateway (192.168.42.129)  0.729 ms  0.829 ms  0.922 mstcpdump: listening on enx0207646b3637  
  
 2  * * *  
 3  255.0.0.1 (255.0.0.1)  67.427 ms  68.297 ms  68.819 ms  
 4  255.0.0.2 (255.0.0.2)  67.964 ms  67.621 ms  68.389 ms  
 5  * * *  
 6  255.0.0.4 (255.0.0.4)  66.334 ms  63.343 ms  63.366 ms  
 7  172.26.100.243 (172.26.100.243)  62.935 ms  172.26.100.242 (172.26.100.242)  35.445 ms  3  
 8  192.168.188.52 (192.168.188.52)  43.909 ms  192.168.188.48 (192.168.188.48)  43.917 ms  4  
 9  * * *  
10  * * *  
11  * * *  
12  * 192.178.110.129 (192.178.110.129)  37.376 ms *  
13  192.178.86.240 (192.178.86.240)  47.810 ms  142.250.214.106 (142.250.214.106)  47.222 ms  
14  142.250.208.221 (142.250.208.221)  47.485 ms  bom12s12-in-f14.1e100.net (142.250.183.78)  
82 packets captured  
94 packets received by filter  
0 packets dropped by kernel
```

We can do a similar execution of `tracert` (on Windows) :

```

C:\Users\ASUS> nslookup google.com
Server: dns-prod.skylus.lan
Address: 10.0.136.7

Non-authoritative answer:
Name: google.com
Addresses: 2404:6800:4009:80c::200e
142.251.220.14

C:\Users\ASUS>tracert 142.251.220.14

Tracing route to hkg07s49-in-f14.1e100.net [142.251.220.14]
over a maximum of 30 hops:

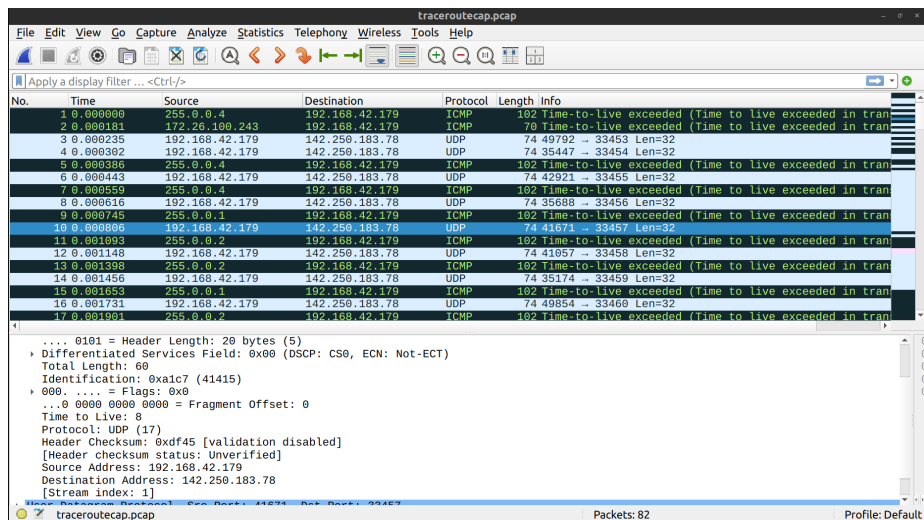
 1      2 ms      1 ms      1 ms  10.7.0.5
 2     43 ms      8 ms      1 ms  172.16.4.7
 3      5 ms      3 ms      3 ms  14.139.98.1
 4     16 ms      1 ms     15 ms  10.117.81.253
 5     12 ms     12 ms     12 ms  10.154.8.137
 6     13 ms     10 ms     12 ms  10.255.239.170
 7     10 ms     52 ms     76 ms  10.152.7.214
 8     61 ms     27 ms     46 ms  72.14.204.62
 9     13 ms     13 ms     13 ms  142.251.76.33
10     18 ms     14 ms     14 ms  142.251.64.13
11    236 ms    162 ms     68 ms  hkg07s49-in-f14.1e100.net [142.251.220.14]

Trace complete.

C:\Users\ASUS>

```

Now, the PCAP file (for Linux `traceroute` execution) can be analysed in WireShark.



Notice that the size of the frame is 60 bytes, just as displayed by `traceroute`. The initial few UDP and ICMP packets are for the Ethernet connection (interface `enx0207646b3637`) through USB to my smartphone.

No.	Time	Source	Destination	Protocol	Length	Info
21	0.035640	172.26.100.242	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
22	0.035716	172.26.100.242	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
23	0.035900	192.168.42.179	142.250.183.78	UDP	74	48621 → 33463 Len=32
24	0.036036	192.168.42.179	142.250.183.78	UDP	74	41883 → 33464 Len=32
25	0.044344	192.168.188.52	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
26	0.044525	192.168.188.48	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
27	0.044623	192.168.42.179	142.250.183.78	UDP	74	42629 → 33465 Len=32
28	0.044723	192.168.188.48	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
29	0.044783	192.168.42.179	142.250.183.78	UDP	74	36780 → 33466 Len=32
30	0.044869	192.168.42.179	142.250.183.78	UDP	74	48023 → 33467 Len=32
31	0.619379	192.168.42.179	192.168.42.129	DNS	82	Standard query 0xb5c5 PTR 1.0.0.255.in-addr.arpa
32	0.620106	192.168.42.129	192.168.42.179	DNS	82	Standard query response 0xb5c5 No such name PTR 1.0.0.255.in-addr.arpa
33	0.626236	192.168.42.179	192.168.42.129	DNS	82	Standard query 0xe2a3 PTR 2.0.0.255.in-addr.arpa
34	0.627699	192.168.42.129	192.168.42.179	DNS	82	Standard query response 0xe2a3 No such name PTR 2.0.0.255.in-addr.arpa
35	0.629055	192.168.42.179	192.168.42.129	DNS	82	Standard query 0x39b7 PTR 4.0.0.255.in-addr.arpa
36	0.629485	192.168.42.129	192.168.42.179	DNS	82	Standard query response 0x39b7 No such name PTR 4.0.0.255.in-addr.arpa
37	0.630562	192.168.42.179	192.168.42.129	DNS	87	Standard query 0x2fc4 PTR 243.160.26.172.in-addr.arpa

Frame 23: 74 bytes on wire (592 bits), 74 bytes captured (592 bits)
 Ethernet II, Src: MS-NLB-PhysServer-07_64:6b:36:37 (02:07:64:6b:36:37), Dst: ae:3d:6c:26:94:dc (ae:3d:6c:26:94:dc)
 Internet Protocol Version 4, Src: 192.168.42.179, Dst: 142.250.183.78
 User Datagram Protocol, Src Port: 48621, Dst Port: 33463
 Data (32 bytes)

Since each packet is being sent through this, it also follows the Ethernet II protocol.

Also notice that although we tried to remove the DNS lookups for `google.com`, there are still some other DNS lookups. Since the probe packets do not follow the DNS message format, we can simply remove the DNS related packets and only focus on the UDP and ICMP/ICMPv6 packets. This can be done using the filter `not dns` in Wireshark.

No.	Time	Source	Destination	Protocol	Length	Info
51	0.724347	192.178.110.129	192.168.42.179	ICMP	110	Time-to-live exceeded (Time to live exceeded in tran
53	0.734393	142.250.214.106	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
54	0.734549	142.251.64.10	192.168.42.179	ICMP	110	Time-to-live exceeded (Time to live exceeded in tran
55	0.734788	142.250.208.221	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
56	0.734923	192.178.86.240	192.168.42.179	ICMP	110	Time-to-live exceeded (Time to live exceeded in tran
57	0.864456	fe80::ac3d:6c:ff:fe2...	2409:4080:8d0d:3b84...	ICMPv6	86	Neighbor Solicitation for 2409:4080:8d0d:3b84:e137:8...
58	0.864540	2409:4080:8d0d:3b84...	fe80::ac3d:6c:ff:fe2...	ICMPv6	78	Neighbor Advertisement 2409:4080:8d0d:3b84:e137:8...
60	0.942777	192.168.42.179	142.250.183.78	UDP	74	48011 → 33474 Len=32
61	0.942849	192.168.42.179	142.250.183.78	UDP	74	38569 → 33475 Len=32
62	0.942890	192.168.42.179	142.250.183.78	UDP	74	46251 → 33476 Len=32
63	0.942933	192.168.42.179	142.250.183.78	UDP	74	52684 → 33477 Len=32
64	0.942972	192.168.42.179	142.250.183.78	UDP	74	50828 → 33478 Len=32
65	0.943011	192.168.42.179	142.250.183.78	UDP	74	50663 → 33479 Len=32
66	0.943049	192.168.42.179	142.250.183.78	UDP	74	53498 → 33480 Len=32
67	0.943087	192.168.42.179	142.250.183.78	UDP	74	36865 → 33481 Len=32
68	0.943127	192.168.42.179	142.250.183.78	UDP	74	46996 → 33482 Len=32
69	0.943170	192.168.42.179	142.250.183.78	UDP	74	34435 → 33483 Len=32
70	0.943210	192.168.42.179	142.250.183.78	UDP	74	45057 → 33484 Len=32
72	0.984605	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
73	0.985600	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
74	0.986691	192.178.110.209	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
75	0.986926	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
76	0.987073	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
77	0.987381	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
78	0.987466	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
79	0.987648	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
80	0.987836	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
81	0.988003	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
82	0.988182	192.178.110.111	192.168.42.179	ICMP	110	Time-to-live exceeded (Time to live exceeded in tran

Frame 30: 74 bytes on wire (592 bits), 74 bytes captured (592 bits)
 Domain Name System Protocol

Even the ICMPv6 packets are not relevant since they are not signifying either TTL reaching 0 or the port being inactive on the destination.

So, we can filter them out too using `(not dns)` and `(not icmpv6)` as the filter.

No.	Time	Source	Destination	Protocol	Length	Info
20	0.002279	192.168.42.179	142.250.183.78	UDP	74	45239 → 33462 Len=32
21	0.035640	172.26.100.242	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
22	0.035716	172.26.100.242	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
23	0.035980	192.168.42.179	142.250.183.78	UDP	74	48621 → 33463 Len=32
24	0.036036	192.168.42.179	142.250.183.78	UDP	74	41883 → 33464 Len=32
25	0.044344	192.168.188.52	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
26	0.044525	192.168.188.48	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
27	0.044623	192.168.42.179	142.250.183.78	UDP	74	42629 → 33465 Len=32
28	0.044726	192.168.188.48	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
29	0.044783	192.168.42.179	142.250.183.78	UDP	74	36780 → 33466 Len=32
30	0.044869	192.168.42.179	142.250.183.78	UDP	74	48923 → 33467 Len=32
45	0.680992	192.168.42.179	142.250.183.78	UDP	74	48698 → 33468 Len=32
46	0.687069	192.168.42.179	142.250.183.78	UDP	74	45963 → 33469 Len=32
47	0.687123	192.168.42.179	142.250.183.78	UDP	74	51856 → 33470 Len=32
48	0.687181	192.168.42.179	142.250.183.78	UDP	74	53081 → 33471 Len=32
49	0.687231	192.168.42.179	142.250.183.78	UDP	74	34558 → 33472 Len=32
50	0.687295	192.168.42.179	142.250.183.78	UDP	74	35343 → 33473 Len=32
51	0.724347	192.178.110.120	192.168.42.179	ICMP	110	Time-to-live exceeded (Time to live exceeded in tran
53	0.734393	142.250.214.106	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
54	0.734549	142.251.64.10	192.168.42.179	ICMP	110	Time-to-live exceeded (Time to live exceeded in tran
55	0.734768	142.250.208.221	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
56	0.734923	192.178.86.240	192.168.42.179	ICMP	110	Time-to-live exceeded (Time to live exceeded in tran
60	0.942777	192.168.42.179	142.250.183.78	UDP	74	48011 → 33474 Len=32
61	0.942849	192.168.42.179	142.250.183.78	UDP	74	38569 → 33475 Len=32
62	0.942890	192.168.42.179	142.250.183.78	UDP	74	46251 → 33476 Len=32
63	0.942933	192.168.42.179	142.250.183.78	UDP	74	52684 → 33477 Len=32
64	0.942972	192.168.42.179	142.250.183.78	UDP	74	50828 → 33478 Len=32
65	0.943011	192.168.42.179	142.250.183.78	UDP	74	50663 → 33479 Len=32
66	0.943049	192.168.42.179	142.250.183.78	UDP	74	53498 → 33480 Len=32

This leaves us with only 63 packets out of the original 82 captured.

Questions

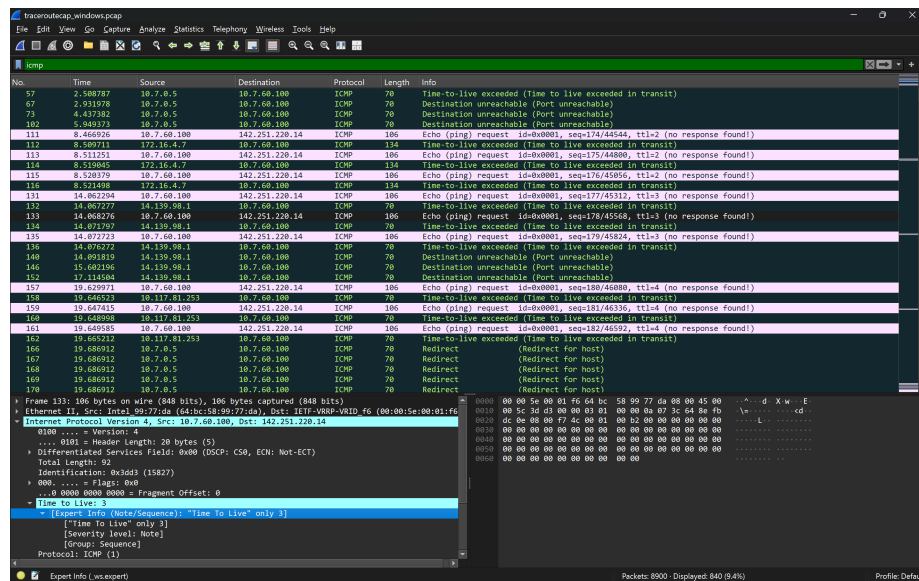
Question 1

What protocol does Windows `tracert` use by default, and what protocol does Linux `traceroute` use by default?

On Linux, the protocol that is used by default for the probe packets is UDP. This is evident from the screenshots of the packets analysed on WireShark for Linux.

But this is only a default settings. We can explicitly state the protocol to use, using `-I` for ICMP, `-T` for TCP SYN packets as probes and `-U` for UDP. There are other protocols too, such as layer-four-TCP (`1ft`) and raw IP datagrams.

For Windows `tracert`, the protocol is ICMP Echo/Ping request by default. For example, in the capture `tracertecap_windows.pcap` :



From the page for `tracert` on Microsoft Learn :

`tracert`- this diagnostic tool determines the path taken to a destination by sending Internet Control Message Protocol (ICMP) echo Request or ICMPv6 messages to the destination with incrementally increasing time to live (TTL) field values

Question 2

Some hops in your traceroute output may show ***. Provide at least two reasons why a router might not reply.

According to the manual page for `traceroute` :

If there is no response within a certain timeout, an "*" (asterisk) is printed for that probe.

This timeout could happen due to various reasons, such as

- the probe packet being lost (in case of UDP probes for example)
 - the Round-Trip-Time being too large for that router/host.
 - A firewall preventing the host from responding to UDP packets. For example, one such firewall exists for `dominos.com`. This is talked about in detail in the answer for Q5.
 - A router not sending ICMP (time to live exceeded) responses for messages that reach TTL value of 0 when being processed by it. In such a case, the probes for the corresponding initial TTL will not get a response, but ones after them may still get responses since those responses will be generated by different routers or the destination end-system.
-

Question 3

In Linux traceroute, which field in the probe packets changes between successive probes sent to the destination?

The source and destination port numbers. Consider these highlighted packets :

46	0.687069	192.168.42.179	142.250.183.78	UDP	74	45063	→	33469	Len=32
47	0.687123	192.168.42.179	142.250.183.78	UDP	74	51856	→	33470	Len=32
48	0.687181	192.168.42.179	142.250.183.78	UDP	74	53081	→	33471	Len=32
49	0.687231	192.168.42.179	142.250.183.78	UDP	74	34558	→	33472	Len=32
50	0.687295	192.168.42.179	142.250.183.78	UDP	74	35343	→	33473	Len=32

In the first highlighted packet, 51856 is the value of the source port number field in the UDP datagram/segment and 33470 is the value for destination port number. All these three highlighted packets had a TTL of 13 as you can see from these three screenshots :

Question 4

At the final hop, how is the response different compared to the intermediate hop?

In the final hop, the response indicates an unreachable port rather than the TTL value reaching 0 (for the packet, as it hops). This is explicitly stated in the manual page for `traceroute` :

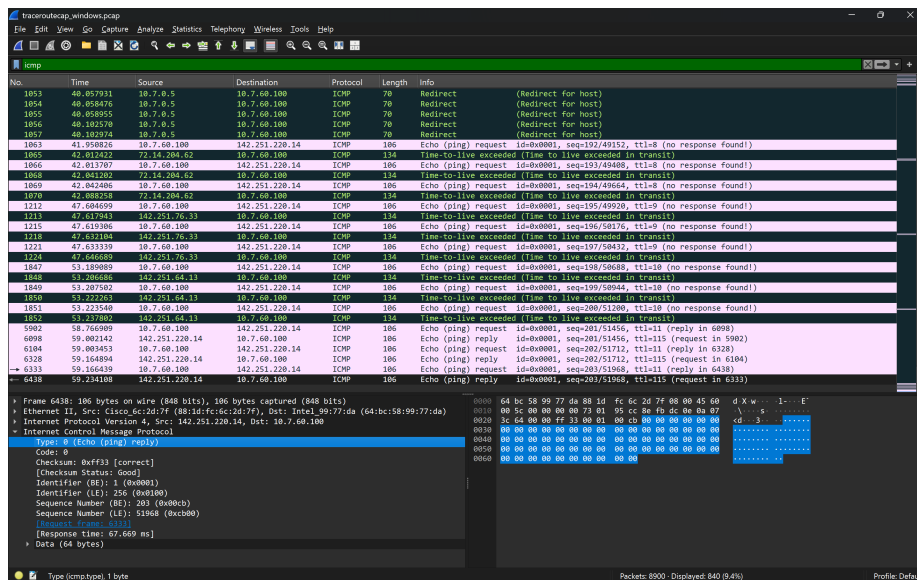
We start our probes with a ttl of one and increase by one until we get an ICMP “port unreachable” (or TCP reset)

For the packets I captured, this shows for the packets at the very end :

64	0.942972	192.168.42.179	142.250.183.78	UDP	74	50828 → 33478 Len=32
65	0.943011	192.168.42.179	142.250.183.78	UDP	74	50863 → 33479 Len=32
66	0.943040	192.168.42.179	142.250.183.78	UDP	74	53498 → 33480 Len=32
67	0.943087	192.168.42.179	142.250.183.78	UDP	74	36865 → 33481 Len=32
68	0.943127	192.168.42.179	142.250.183.78	UDP	74	46996 → 33482 Len=32
69	0.943170	192.168.42.179	142.250.183.78	UDP	74	34435 → 33483 Len=32
70	0.943210	192.168.42.179	142.250.183.78	UDP	74	45057 → 33484 Len=32
72	0.984665	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
73	0.996560	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
74	0.996691	192.178.110.209	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in transit)
75	0.996928	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
76	0.997073	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
77	0.997301	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
78	0.997466	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
79	0.997648	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
80	0.997830	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
81	0.998003	142.250.183.78	192.168.42.179	ICMP	70	Destination unreachable (Port unreachable)
82	0.998182	192.178.110.111	192.168.42.179	ICMP	110	Time-to-live exceeded (Time to live exceeded in transit)

On the other hand, for an intermediate hop, the expected response (for Linux) is an ICMP packet saying “time to live exceeded”.

On Windows, the expected type of response from the destination end-system is an ICMP echo reply.



The screenshot shows the Wireshark interface with the following details:

- Packet List:** Shows a list of ICMP Echo (ping) requests and replies. The first packet is an Echo (ping) request from 10.7.68.100 to 142.251.228.14. The last packet is an Echo (ping) reply from 142.251.228.14 to 10.7.68.100.
- Packet Details:** Shows the structure of an ICMP Echo (ping) request. The fields are: Type: 8 (Echo), Code: 0, Checksum: 0x0f03 [correct], Identifier (NE): 1 (0x0001), Identifier (LE): 256 (0x0100), Sequence Number (BE): 203 (0x00cb), Sequence Number (LE): 51968 (0xc000).
- Packet Bytes:** Shows the raw data of the ICMP Echo (ping) request. The data is: 0000 64 bc 58 99 77 da 88 1d fc 6c 2d 7f 00 00 45 60 d X w . . 1 . . E .

Question 5

Suppose a firewall blocks UDP traffic but allows ICMP — how would this affect the results of Linux `traceroute` vs. Windows `tracert`?

Many ISPs and CDNs (like Akamai, which serves `dominos.com`) configure intermediate routers to ignore ICMP Echo Requests or rate-limit ICMP responses. This prevents their infrastructure from being overloaded by diagnostic tools like `traceroute/ping`. Let's take the example of `dominos.com`.

```
$ sudo traceroute dominos.com -n 16
traceroute to dominos.com (23.9.119.229), 16 hops max, 60 byte packets
 1 10.240.0.2 (10.240.0.2) 3.667 ms 3.585 ms 3.539 ms
 2 10.3.0.29 (10.3.0.29) 5.815 ms 5.758 ms 5.711 ms
 3 10.3.0.5 (10.3.0.5) 5.659 ms 5.613 ms 5.562 ms
 4 172.16.4.7 (172.16.4.7) 3.208 ms 3.152 ms 3.094 ms
 5 14.139.98.1 (14.139.98.1) 5.605 ms 5.550 ms 5.497 ms
 6 10.117.81.253 (10.117.81.253) 5.170 ms 6.578 ms 6.500 ms
 7 * * *
 8 * * *
 9 10.255.222.33 (10.255.222.33) 84.069 ms 83.937 ms 10.255.222.33 (10.255.222.33) 83.876 ms
10 115.247.100.29 (115.247.100.29) 83.824 ms 83.766 ms 83.711 ms
11 * * *
12 121.240.252.1.static-hyderabad.vsnl.net.in (121.240.252.1) 83.365 ms * 30.098 ms
13 * * *
14 121.244.3.222.static-mumbai.vsnl.net.in (121.244.3.222) 43.713 ms 43.668 ms 29.583 ms
15 * * *
16 * * *
$ sudo traceroute dominos.com -n 16 -I
traceroute to dominos.com (23.9.119.229), 16 hops max, 60 byte packets
 1 10.240.0.2 (10.240.0.2) 2.614 ms 2.586 ms 2.581 ms
 2 10.3.0.29 (10.3.0.29) 4.097 ms 4.093 ms 4.087 ms
 3 10.3.0.5 (10.3.0.5) 4.082 ms 4.074 ms 4.069 ms
 4 172.16.4.7 (172.16.4.7) 2.411 ms 2.404 ms 2.396 ms
 5 14.139.98.1 (14.139.98.1) 4.968 ms 4.962 ms 4.956 ms
 6 10.117.81.253 (10.117.81.253) 3.987 ms 6.429 ms 6.409 ms
 7 * * *
 8 * * *
 9 10.255.222.33 (10.255.222.33) 89.719 ms 89.693 ms 89.674 ms
10 115.247.100.29 (115.247.100.29) 89.659 ms 89.647 ms 89.619 ms
11 * * *
12 121.240.252.1.static-hyderabad.vsnl.net.in (121.240.252.1) 89.510 ms 43.614 ms 43.575 ms
13 * * *
14 121.244.3.222.static-mumbai.vsnl.net.in (121.244.3.222) 101.342 ms 101.324 ms 57.720 ms
15 a23-9-119-229.deploy.static.akamaitechnologies.com (23.9.119.229) 57.673 ms 102.124 ms 102.083 ms
```

As you can see from this execution of `traceroute`, the UDP probes are unable to reach the final destination (as indicated by the astericks for TTL values 15 and 16), but the ICMP probes are.

Since `tracert` on Windows uses ICMP probes by default, the probes aren't dropped pre-maturely and are eventually able to reach the destination (as TTL is increased).

```
PS C:\Users\ASUS> tracert dominos.com
```

Tracing route to a23-9-119-229.deploy.static.akamaitechnologies.com [23.9.119.229] over a maximum of 30 hops:

1	3 ms	2 ms	2 ms	10.27.13.154
2	29 ms	5 ms	6 ms	10.7.0.5
3	22 ms	10 ms	10 ms	172.16.4.7
4	38 ms	7 ms	8 ms	14.139.98.1
5	43 ms	6 ms	6 ms	10.117.81.253

6	*	*	*	Request timed out.
7	*	*	*	Request timed out.
8	47 ms	34 ms	39 ms	10.255.222.33
9	62 ms	36 ms	35 ms	115.247.100.29
10	*	*	*	Request timed out.
11	61 ms	35 ms	40 ms	121.240.252.1.static-hyderabad.vsnl.net.in [121.240.252.1]
12	*	*	*	Request timed out.
13	167 ms	84 ms	42 ms	121.244.3.222.static-mumbai.vsnl.net.in [121.244.3.222]
14	70 ms	32 ms	32 ms	a23-9-119-229.deploy.static.akamaitechnologies.com [23.9.119.229]

Trace complete.

Does using IP address for traceroute make a difference ?

Now, since doing this for one of the servers that `google.com` may resolve to may not be satisfactory for whoever is going to grade this assignment, I will *also* capture the packets for `traceroute google.com` and show the differences.

The script will be the same as before, but with `www.google.com` instead of `142.251.220.14`.

The output of `traceroute` is :

```
tcpdump: listening on enx0207646b3637, link-type EN10MB (Ethernet), snapshot length 262144 bytes
traceroute to www.google.com (142.250.192.4), 30 hops max, 60 byte packets
 1 _gateway (192.168.42.129) 0.671 ms 0.767 ms 0.851 ms
 2 * * *
 3 255.0.0.1 (255.0.0.1) 50.453 ms 51.420 ms 52.011 ms
 4 255.0.0.2 (255.0.0.2) 51.062 ms 51.730 ms 52.224 ms
 5 * * *
 6 255.0.0.4 (255.0.0.4) 49.586 ms 46.914 ms 46.343 ms
 7 172.26.100.242 (172.26.100.242) 47.384 ms 172.26.100.243 (172.26.100.243) 28.715 ms 172.26.100.242 (172.26.100.242) 28.697 ms
 8 192.168.188.50 (192.168.188.50) 39.180 ms 192.168.188.48 (192.168.188.48) 39.082 ms 38.972 ms
 9 * * *
10 * * *
11 173.194.121.8 (173.194.121.8) 49.866 ms * *
12 192.178.110.227 (192.178.110.227) 39.026 ms * *
13 142.250.208.223 (142.250.208.223) 61.612 ms 216.239.46.136 (216.239.46.136) 72.893 ms 216.239.50.166 (216.239.50.166) 73.304 ms
14 142.250.208.223 (142.250.208.223) 73.019 ms 192.178.110.248 (192.178.110.248) 73.582 ms 108.170.231.79 (108.170.231.79) 73.347 ms
15 bom12s14-in-f4.1e100.net (142.250.192.4) 37.230 ms 40.752 ms 192.178.110.205 (192.178.110.205) 40.848 ms
129 packets captured
130 packets received by filter
0 packets dropped by kernel
^C
```

For better visibility, this is the output in text :

```
tcpdump: listening on enx0207646b3637, link-type EN10MB (Ethernet), snapshot length 262144 bytes
traceroute to www.google.com (142.250.192.4), 30 hops max, 60 byte packets
 1 _gateway (192.168.42.129) 0.671 ms 0.767 ms 0.851 ms
 2 * * *
 3 255.0.0.1 (255.0.0.1) 50.453 ms 51.420 ms 52.011 ms
 4 255.0.0.2 (255.0.0.2) 51.062 ms 51.730 ms 52.224 ms
 5 * * *
 6 255.0.0.4 (255.0.0.4) 49.586 ms 46.914 ms 46.343 ms
 7 172.26.100.242 (172.26.100.242) 47.384 ms 172.26.100.243 (172.26.100.243) 28.715 ms 172.26.100.242 (172.26.100.242) 28.697 ms
 8 192.168.188.50 (192.168.188.50) 39.180 ms 192.168.188.48 (192.168.188.48) 39.082 ms 38.972 ms
 9 * * *
10 * * *
11 173.194.121.8 (173.194.121.8) 49.866 ms * *
12 192.178.110.227 (192.178.110.227) 39.026 ms * *
13 142.250.208.223 (142.250.208.223) 61.612 ms 216.239.46.136 (216.239.46.136) 72.893 ms 216.239.50.166 (216.239.50.166) 73.304 ms
14 142.250.208.223 (142.250.208.223) 73.019 ms 192.178.110.248 (192.178.110.248) 73.582 ms 108.170.231.79 (108.170.231.79) 73.347 ms
15 bom12s14-in-f4.1e100.net (142.250.192.4) 37.230 ms 40.752 ms 192.178.110.205 (192.178.110.205) 40.848 ms
129 packets captured
130 packets received by filter
0 packets dropped by kernel
```

When analysed in WireShark, without any filters, we get this :

Apply a display filter ... <Ctrl-/>

No.	Time	Source	Destination	Protocol	Length	Info
1	0.009080	192.168.42.129	192.168.42.179	DNS	99	Standard query response 0xb3c9 A www.google.com A 142.250.192.4
2	0.019268	192.168.42.129	192.168.42.179	DNS	102	Standard query response 0xb74c AAAA www.google.com A 142.250.192.4
3	0.010908	192.168.42.179	142.250.192.4	UDP	74	41449 → 33435 Len=32
4	0.011030	192.168.42.179	142.250.192.4	UDP	74	36342 → 33435 Len=32
5	0.011072	192.168.42.179	142.250.192.4	UDP	74	51180 → 33436 Len=32
6	0.011107	192.168.42.179	142.250.192.4	UDP	74	55009 → 33437 Len=32

↳ Frame 1: 99 bytes on wire (728 bits), 99 bytes captured (728 bits)
 ↳ Ethernet II, Src: 16:63:b9:ed:a5:af (16:63:b9:ed:a5:af), Dst: MS-MLB-PhysServer-07_64:6b:36:37 (02:07:64:6b:36:37)
 ↳ Internet Protocol Version 4, Src: 192.168.42.129, Dst: 192.168.42.179
 ↳ User Datagram Protocol, Src Port: 53, Dst Port: 49016
 ↳ Domain Name System (response)
 Transaction ID: 0xb3c9
 Flags: 0x1800 Standard query response, No error
 Questions: 1
 Answer RRs: 1
 Authority RRs: 0
 Additional RRs: 0
 Queries
 ↳ www.google.com: type A, class IN
 Answers
 ↳ www.google.com: type A, class IN, addr 142.250.192.4
 [Unsolicited: True]

As you can see, before sending probes to `www.google.com`, we first need to do a DNS lookup.

The IP address returned is `142.250.192.4` which is different from what we got earlier from `nslookup`. This is because of load distribution/balancing over the servers.

There were also no ICMPv6 responses when using `google.com` directly.

Apart from that, there were no major differences.

(not dns)

No.	Time	Source	Destination	Protocol	Length	Info
16	0.011421	192.168.42.179	142.250.192.4	UDP	74	43951 → 33447 Len=32
17	0.011448	192.168.42.179	142.250.192.4	UDP	74	51704 → 33448 Len=32
18	0.011476	192.168.42.179	142.250.192.4	UDP	74	43868 → 33449 Len=32
19	0.011587	192.168.42.129	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
20	0.011768	192.168.42.129	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
21	0.011913	192.168.42.129	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
24	0.014410	192.168.42.179	142.250.192.4	UDP	74	24594 → 33450 Len=32
25	0.014540	192.168.42.179	142.250.192.4	UDP	74	33240 → 33451 Len=32
26	0.014586	192.168.42.179	142.250.192.4	UDP	74	50922 → 33452 Len=32
27	0.060870	255.0.0.4	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
28	0.061055	255.0.0.4	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
29	0.061167	192.168.42.179	142.250.192.4	UDP	74	56110 → 33453 Len=32
30	0.061306	255.0.0.1	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
31	0.061347	192.168.42.179	142.250.192.4	UDP	74	47617 → 33454 Len=32
32	0.061476	192.168.42.179	142.250.192.4	UDP	74	34692 → 33455 Len=32
33	0.061646	255.0.0.1	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
34	0.061788	192.168.42.179	142.250.192.4	UDP	74	46419 → 33456 Len=32
35	0.061950	172.20.100.242	192.168.42.179	ICMP	70	Time-to-live exceeded (Time to live exceeded in tran
36	0.062125	192.168.42.179	142.250.192.4	UDP	74	68047 → 33457 Len=32
37	0.062358	255.0.0.2	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
38	0.062447	192.168.42.179	142.250.192.4	UDP	74	49223 → 33458 Len=32
39	0.062651	255.0.0.1	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
40	0.062732	192.168.42.179	142.250.192.4	UDP	74	30792 → 33459 Len=32
41	0.063050	255.0.0.2	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
42	0.063141	192.168.42.179	142.250.192.4	UDP	74	59856 → 33460 Len=32
43	0.063277	255.0.0.1	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
44	0.063356	192.168.42.179	142.250.192.4	UDP	74	50430 → 33461 Len=32
45	0.063581	255.0.0.2	192.168.42.179	ICMP	102	Time-to-live exceeded (Time to live exceeded in tran
46	0.063664	192.168.42.179	142.250.192.4	UDP	74	35552 → 33462 Len=32